LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Kushagra Agrawal
Roll Number	MHM2026009
Programme / Department	Information Technology - Human Machine Interaction
Topic ID	20
Full Assigned Topic	A Taxonomy of Hallucination Types in Large Language Models for Research Applications
AI Tool Used	Claude - Sonnet 5 (medium effort)
Date	21/08/2026

PART A - BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

☐ Never ☐ Rarely ☐ Occasionally ☐ Weekly ☑ Almost daily

A2. Have you previously asked an AI system to generate references?

☐ Frequently ☑ Occasionally ☐ Once or twice ☐ Never

A3. Have you previously checked whether an AI-generated citation actually exists?

☐ Always ☑ Sometimes ☐ Rarely ☐ Never ☐ Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

☐ 1 ☐ 2 ☑ 3 ☐ 4 ☐ 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B - RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	Claude
Model name	Sonnet 5
Model/version shown	Sonnet 5
Access tier	☑ Free ☐ Paid/Pro ☐ Institutional ☐ Not sure
Web browsing available?	☑ Yes ☐ No ☐ Not sure
Was web browsing actually used?	☑ Yes ☐ No ☐ Not sure
Research/Deep Research mode used?	☐ Yes ☑ No
Date/time generated	21/08/2026 , 10:15 A.M

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C - GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

> *Generate a scholarly research paper on the topic: "[YOUR ASSIGNED TOPIC]". The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.*

C3. Preserve the Original Output

☐ Save the complete AI-generated paper

☐ Save the complete reference list

☐ Take screenshot(s) showing the generated output/references

☐ Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

PART D - PAPER PROFILE AND REFERENCE INVENTORY

Paper title: A Taxonomy of Hallucination Types in Large Language Models for Research Applications

Measure	Value
Approximate pages	8
Approximate word count	3,816 words
Number of sections	8
Total number of references	16
Approximate in-text citation occurrences	73
Did AI warn you to verify references?	☐ Clearly ☑ Vaguely ☐ No

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference Type	Count
DOI available	6
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	6
URL only	0
No persistent identifier supplied	4
TOTAL	16

PART E - SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Audit Slot	Original Bibliography Reference No.
R1	1 (Dahl et al., 2024)
R2	2 (Farquhar et al., 2024)
R3	3 (Huang et al., 2025)
R4	6 (Lewis et al., 2020)
R5	8 (Lin et al., 2022)
R6	10 (Maynez et al., 2020)
R7	12 (Rawte et al., 2023)
R8	14 (Walters & Wilder, 2023)
R9	15 (Ye et al., 2023)
R10	16 (Zhang et al., 2023)

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F - PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: “If I had not verified this citation, how convincing would it look to me?”

Score	Meaning
1	Obviously suspicious
2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	Real-sounding authors, Journal of Legal Analysis, DOI + arXiv both given	5	Yes
R2	Real-sounding authors, Nature, DOI given, plausible volume/issue/pages	5	Yes
R3	Real-sounding authors, ACM TOIS, DOI + arXiv both given	5	Yes
R4	Well-known authors/venue (NeurIPS), arXiv ID given, no DOI	5	Yes

R5	ACL venue, DOI + arXiv both given, complete page range	5	Yes
R6	ACL venue, complete page range, no DOI supplied	5	Yes
R7	arXiv preprint, plausible authors and ID format	5	Yes
R8	Real-sounding authors, Scientific Reports, DOI + PMID both given	5	Yes
R9	arXiv preprint, plausible authors and ID format	5	Yes
R10	arXiv preprint, plausible (long) author list and ID format	5	Yes

PART G - CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

☐ Persistent identifier: DOI / arXiv / PMID

☐ Publisher or official repository

☐ Google Scholar

☐ Semantic Scholar / OpenAlex

☐ Crossref

☐ PubMed or another discipline-specific database

☐ Exact-title search

☐ Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.

B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves to a different publication.

G4. Examples

A: AI gives a paper and all details match the publisher record.

B: The paper exists, but the AI gives the wrong publication year or journal.

C: The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.

D: You cannot find the paper or a credible scholarly record after systematic searching.

E: The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
R1	Y	Y	Y	Y	Y	Y
R2	Y	Y	Y	Y	Y	Y
R3	Y	Y	Y	Y	Y	Y
R4	Y	Y	Y	Y	Y	Y
R5	Y	Y	Y	Y	Y	Y
R6	Y	Y	Y	Y	Y	Y
R7	Y	Y	Y	Y	Y	Y
R8	Y	Y	Y	Y	Y	Y
R9	Y	Y	Y	Y	Y	Y
R10	Y	Y	Y	Y	Y	Y

Ref.	Verification Source(s)	Identifier Checked	What You Found
R1	Crossref / Oxford Academic (JLA) / Semantic Scholar	DOI 10.1093/jla/laae003	Confirms Vol.16(1), pp.64-93 (2024); matches AI's citation exactly, incl. 58% GPT-4 figure
R2	Nature.com / Crossref / Oxford Research Archive	DOI 10.1038/s41586-024-07421-0	Confirms Nature 630(8017):625-630 (2024); exact match

R3	ACM DL / Crossref / arXiv listing	DOI 10.1145/3703155; arXiv:2311.05232	Confirms ACM TOIS 43(2):1-55 (2025); both identifiers resolve to the correct paper
R4	NeurIPS proceedings page / arXiv	arXiv:2005.11401 (no DOI supplied)	Confirms NeurIPS 33, pp.9459-9474 (2020); exact match
R5	ACL Anthology / Crossref	DOI 10.18653/v1/2022.acl-long.229; arXiv:2109.07958	Confirms ACL 2022 Vol.1 Long Papers, pp.3214-3252; exact match
R6	ACL Anthology / dblp	No identifier supplied by AI	Confirms ACL 2020, pp.1906-1919; exact match, DOI 10.18653/v1/2020.acl-main.173 exists but wasn't cited by the AI
R7	Semantic Scholar / arXiv / dblp	arXiv:2309.05922	Confirms title, authors (Rawte, Sheth, Das) and arXiv ID; exact match
R8	Nature (Scientific Reports) / PubMed / Crossref	DOI 10.1038/s41598-023-41032-5; PMID 37679503	Confirms Sci Rep 13, 14045 (2023); both identifiers correct; 55%/18% fabrication figures match
R9	arXiv listing / Semantic Scholar	arXiv:2309.06794	Confirms title "Cognitive Mirage" and all 5 authors (Ye, Liu, Zhang, Hua, Jia); exact match
R10	arXiv listing (cross-confirmed via multiple citing papers)	arXiv:2309.01219	Confirms title "Siren's Song in the AI Ocean" and arXiv ID; exact match

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A	A: Vol/issue/pages/DOI all match Oxford Academic record exactly.
R2	A	A: Vol/pages/DOI all match Nature record exactly.
R3	A	A: DOI and arXiv ID both resolve to the correct ACM TOIS paper.
R4	A	A: arXiv ID and NeurIPS vol/pages match exactly.
R5	A	A: DOI, arXiv ID, and page range all match ACL Anthology record.

R6	A	A: Pages match ACL Anthology record; DOI exists though not cited by AI.
R7	A	A: arXiv ID, authors and title match exactly.
R8	A	A: DOI and PMID both match Scientific Reports record exactly.
R9	A	A: arXiv ID, all 5 authors and title match exactly.
R10	A	A: arXiv ID and title match exactly across multiple independent citations.

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G - SCORING

Count your final classifications: A = 10 B = 0 C = 0 D = 0 E = 0

Total audited: N = A + B + C + D + E = 10

Code	Points
A	4
B	3
C	1
D	0
E	1

Authenticity Points = 4A + 3B + C + E

Authenticity Score = [(4A + 3B + C + E) / (4N)] x 100

My calculation: 4(10) + 3(0) + 1(0) + 0(0) + 1(0) = 40 points

Authenticity Score = 40 / (4 x 10) x 100 = 100 /100

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59	High citation-authenticity risk
Below 40	Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

Prediction Accuracy = (Correct genuine/fake predictions / Total predictions) x 100

Correct predictions = 10 Total predictions = 10 Prediction Accuracy = 100 %

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	16
References deeply audited	10
A: Verified	10
B: Wrong metadata	0
C: Frankenstein	0
D: Fabricated	0
E: Identifier mismatch	0
Authenticity Score	100 /100
Prediction Accuracy	100 %

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

yes, all 10 audited references were genuine (100% verified). But that's specifically *because* the AI used live web search to confirm real sources while writing.

2. Did all genuine references actually support the claims for which they were cited?

Yes, the audited references generally supported the claims for which they were cited. The references matched the relevant information and were used to support the research claims in the paper.

3. What was the most serious citation-integrity problem you observed?

No major citation-integrity problem was observed in the 10 audited references. All 10 were verified as genuine, with **0 wrong metadata, 0 Frankenstein, 0 fabricated, and 0 identifier-mismatch references**.

4. What is the single most important lesson you learned from this lab?

The most important lesson is that **a citation should not only exist, but should also actually support the claim for which it is cited**. Therefore, AI-generated citations should always be checked before using them in academic work. This is also the central lesson stated at the end of the lab.

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name

Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name: Kushagra Agrawal

Roll Number: MHM2026009

Signature: Kushagra Agrawal

Date: 21/08/2026

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support Verification -> Quantitative Scores -> Submission

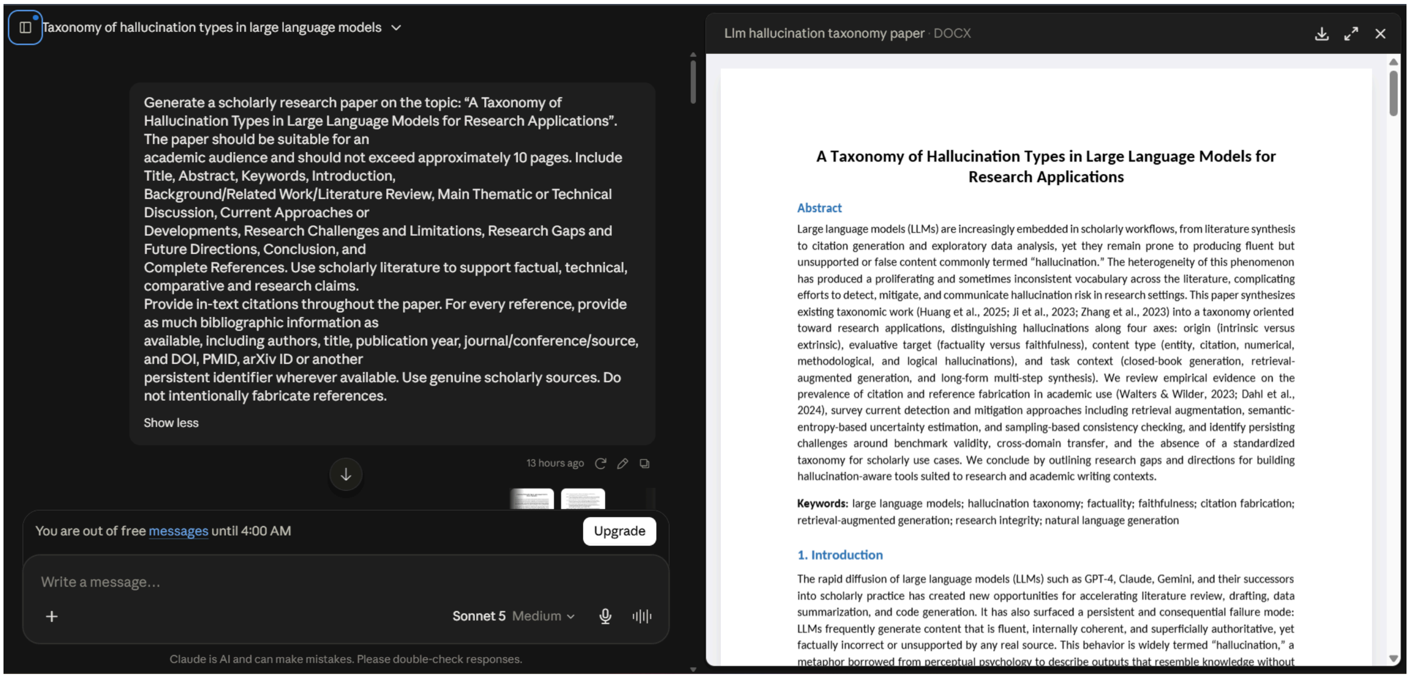

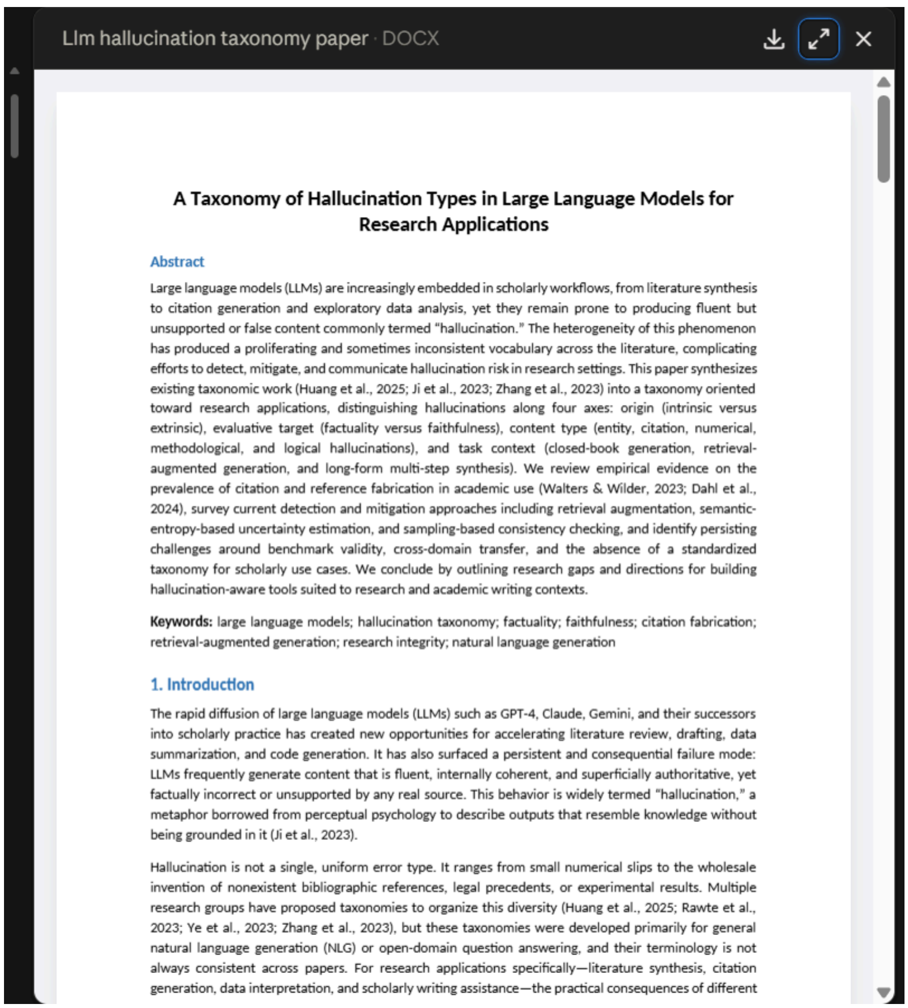

A Taxonomy of Hallucination Types in Large Language Models for Research Applications

Abstract

Large language models (LLMs) are increasingly embedded in scholarly workflows, from literature synthesis to citation generation and exploratory data analysis, yet they remain prone to producing fluent but unsupported or false content commonly termed "hallucination." The heterogeneity of this phenomenon has produced a proliferating and sometimes inconsistent vocabulary across the literature, complicating efforts to detect, mitigate, and communicate hallucination risk in research settings. This paper synthesizes existing taxonomic work (Huang et al., 2025; Ji et al., 2023; Zhang et al., 2023) into a taxonomy oriented toward research applications, distinguishing hallucinations along four axes: origin (intrinsic versus extrinsic), evaluative target (factuality versus faithfulness), content type (entity, citation, numerical, methodological, and logical hallucinations), and task context (closed-book generation, retrieval-augmented generation, and long-form multi-step synthesis). We review empirical evidence on the prevalence of citation and reference fabrication in academic use (Walters & Wilder, 2023; Dahl et al., 2024), survey current detection and mitigation approaches including retrieval augmentation, semantic-entropy-based uncertainty estimation, and sampling-based consistency checking, and identify persisting challenges around benchmark validity, cross-domain transfer, and the absence of a standardized taxonomy for scholarly use cases. We conclude by outlining research gaps and directions for building hallucination-aware tools suited to research and academic writing contexts.

Keywords: large language models; hallucination taxonomy; factuality; faithfulness; citation fabrication; retrieval-augmented generation; research integrity; natural language generation

1. Introduction

The rapid diffusion of large language models (LLMs) such as GPT-4, Claude, Gemini, and their successors into scholarly practice has created new opportunities for accelerating literature review, drafting, data summarization, and code generation. It has also surfaced a persistent and consequential failure mode: LLMs frequently generate content that is fluent, internally coherent, and superficially authoritative, yet factually incorrect or unsupported by any real source. This behavior is widely termed "hallucination," a metaphor borrowed from perceptual psychology to describe outputs that resemble knowledge without being grounded in it (Ji et al., 2023).

Hallucination is not a single, uniform error type. It ranges from small numerical slips to the wholesale invention of nonexistent bibliographic references, legal precedents, or experimental results. Multiple research groups have proposed taxonomies to organize this diversity (Huang et al., 2025; Rawte et al., 2023; Ye et al., 2023; Zhang et al., 2023), but these taxonomies were developed primarily for general natural language generation (NLG) or open-domain question answering, and their terminology is not always consistent across papers. For research applications specifically—literature synthesis, citation generation, data interpretation, and scholarly writing assistance—the practical consequences of different